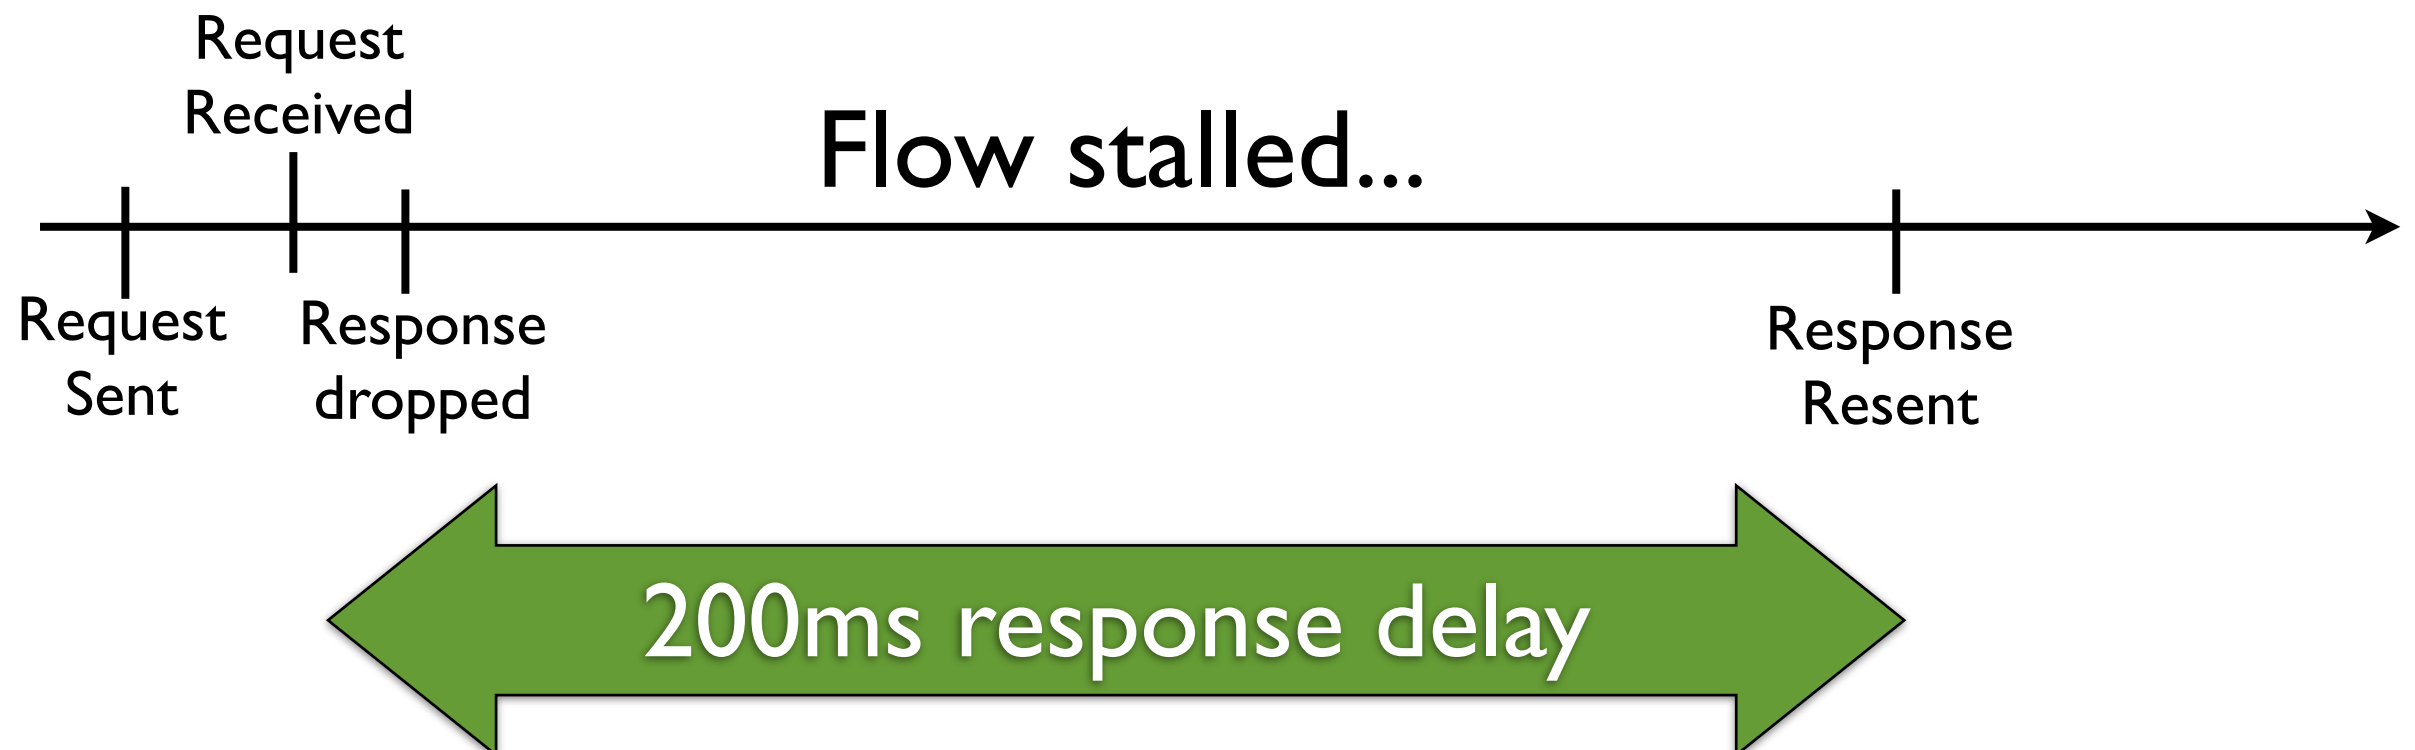



# Safe and Effective Fine-grained TCP Retransmissions for Datacenter Communication

**Vijay Vasudevan**, Amar Phanishayee, Hiral Shah, Elie Krevat  
David Andersen, Greg Ganger, Garth Gibson, Brian Mueller\*

Carnegie Mellon University, \*Panasas Inc.

# Datacenter TCP Request-Response



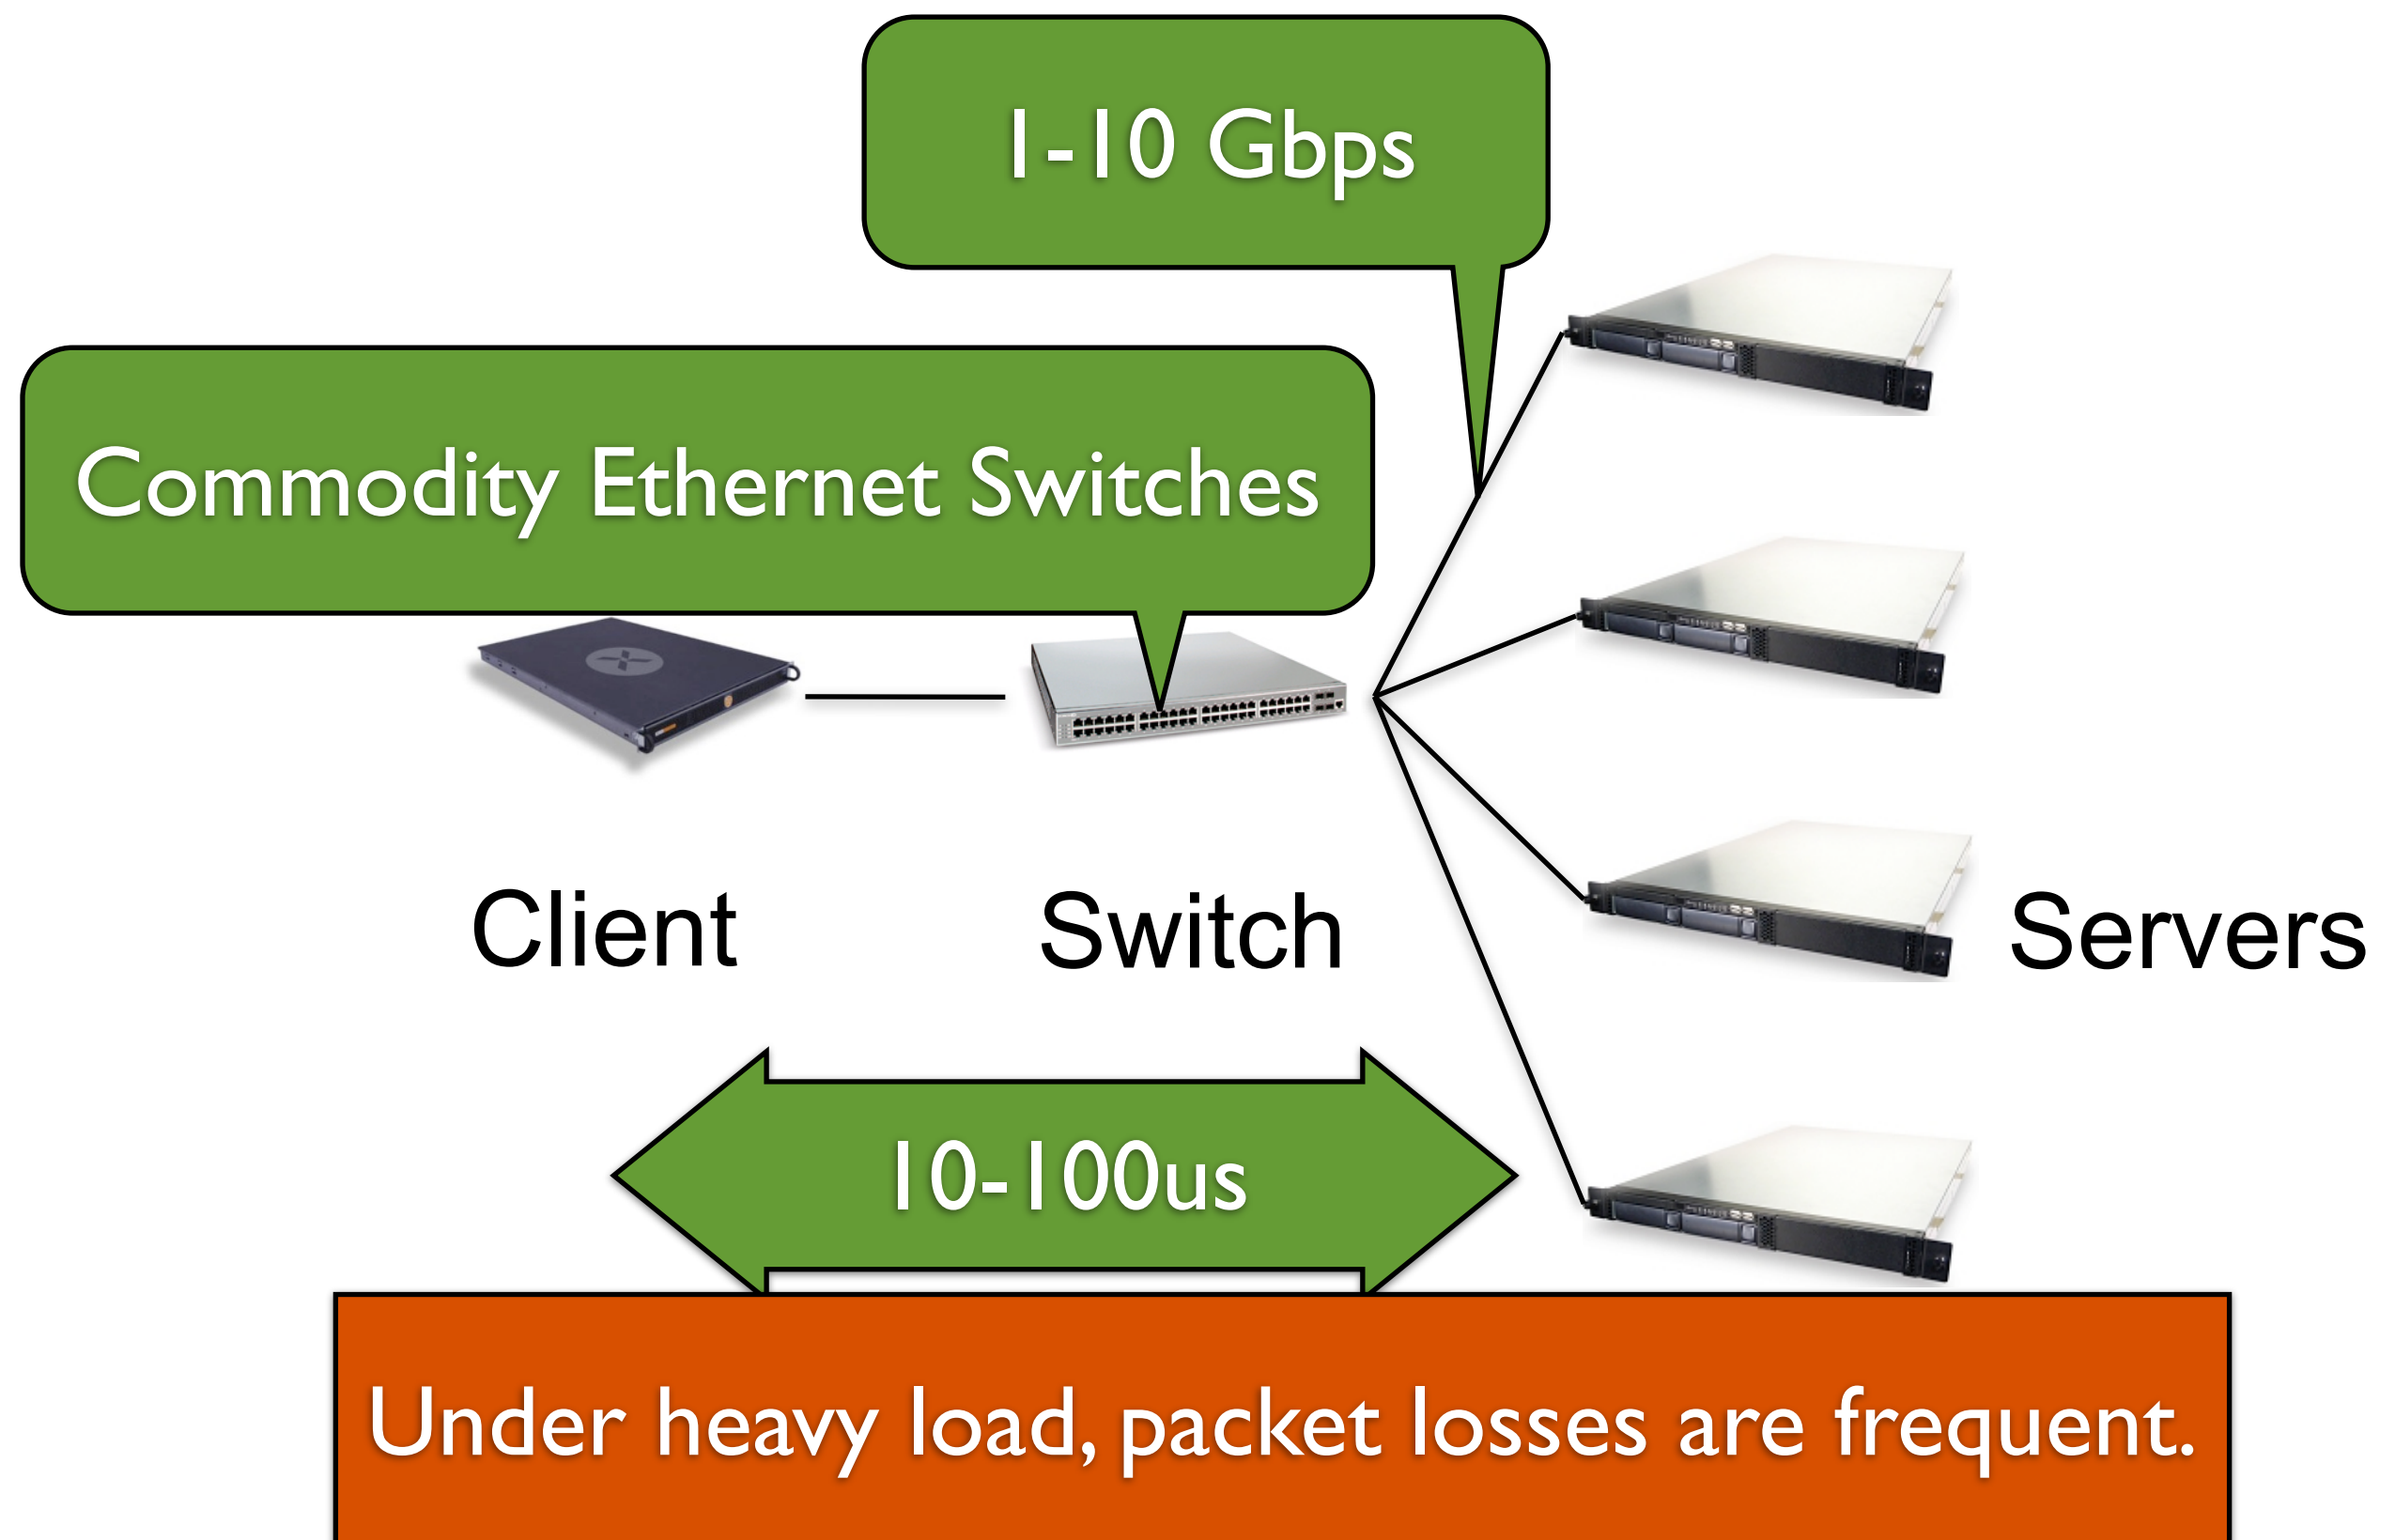
# Applications Sensitive to 200ms TCP Timeouts

- “Drive-bys” affecting single-flow request/response
- Barrier-Sync workloads
  - Parallel cluster filesystems (Incast workloads)
  - Massive multi-server queries (e.g., previous talk)
    - Latency-sensitive, customer-facing

# Main Takeaways


- Problem: 200ms TCP timeouts can cripple datacenter apps
- Solution: Enable microsecond retransmissions
- Can improve datacenter app throughput/latency
- Safe in the wide-area

# The Datacenter Environment



# TCP: Loss Recovery Comparison

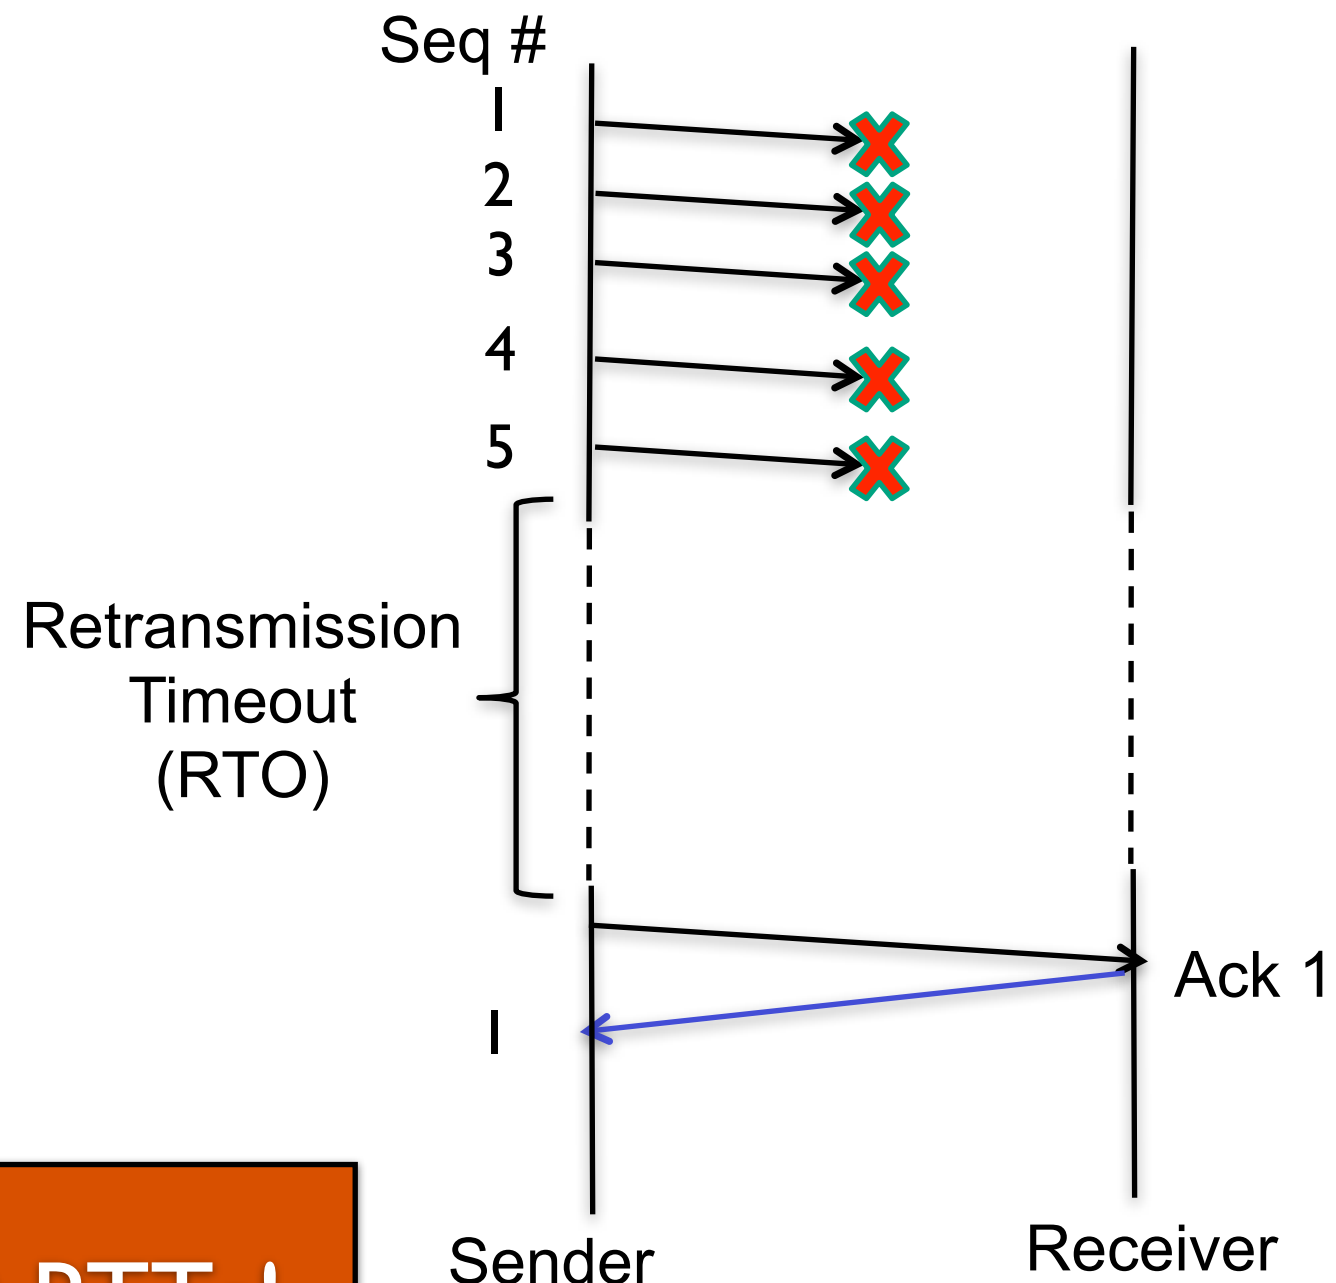
Data-driven recovery is  
**super fast (us)**



minRTO	200.0ms
DC Latency	0.1ms

**1 TCP Timeout lasts 1000 RTTs!**

Timeout driven recovery is  
**painfully slow (ms)**

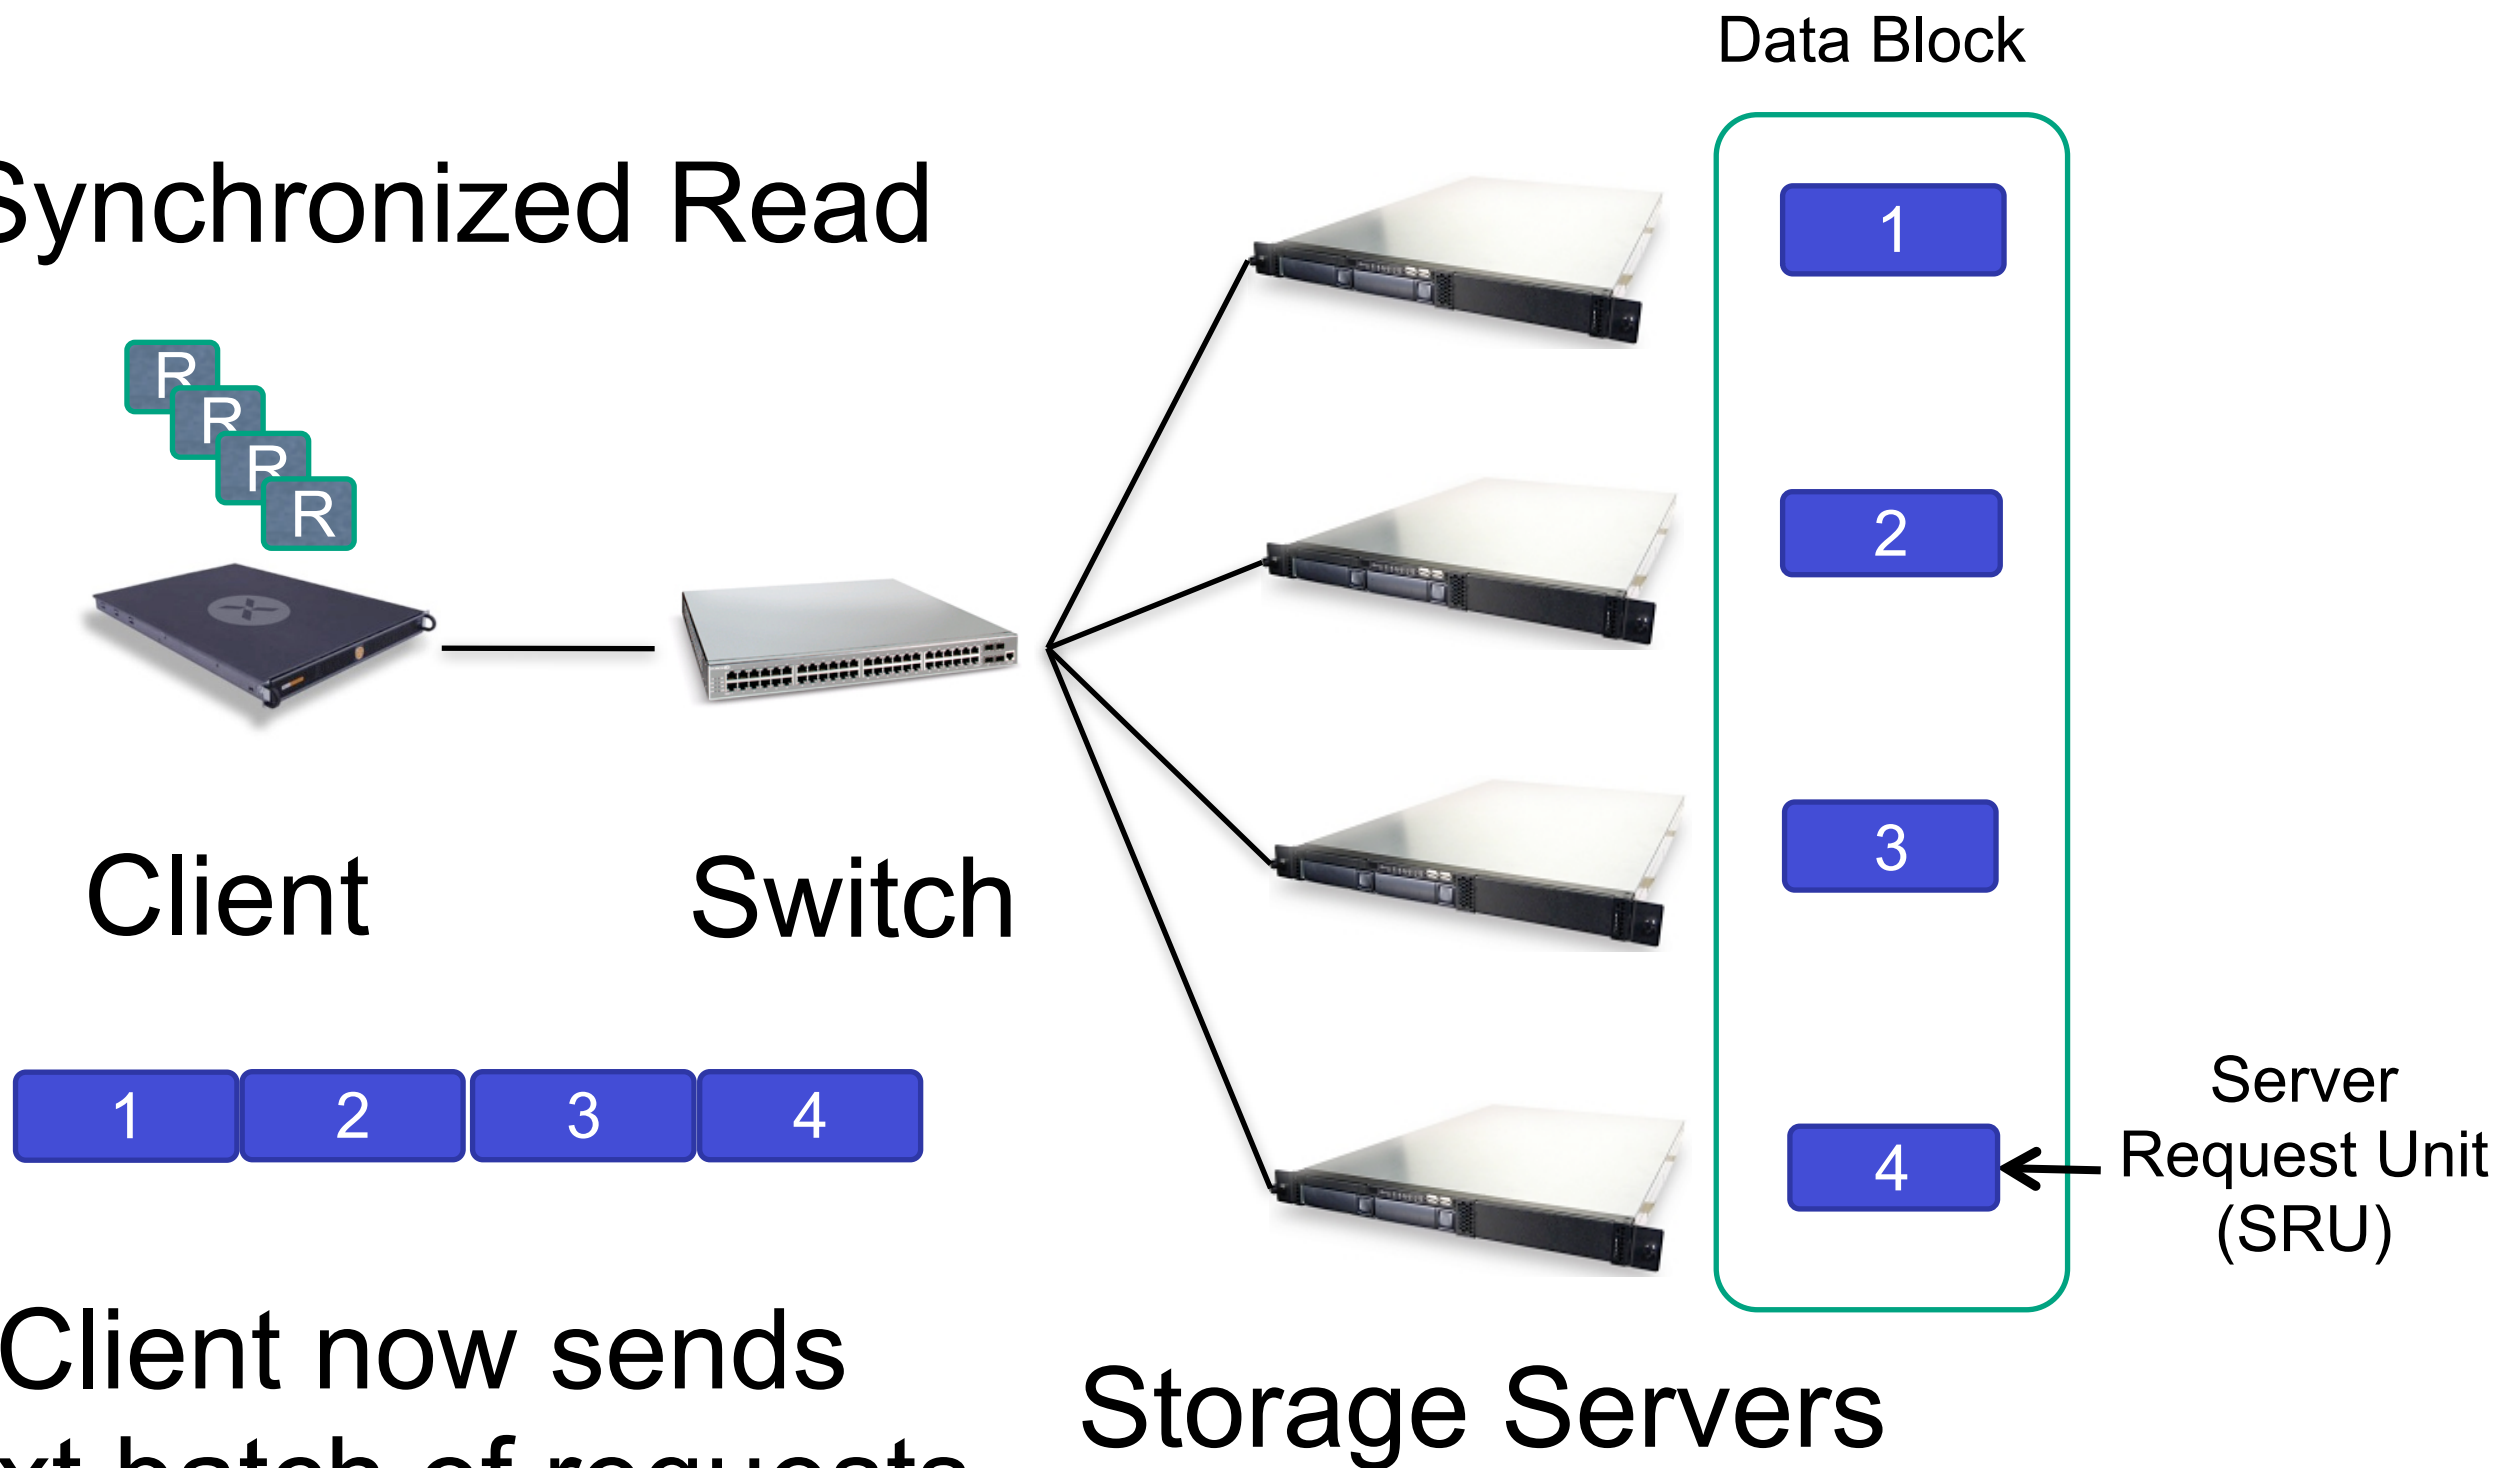


# RTO Estimation and Minimum Bound

- Jacobson's TCP RTO Estimator
  - $RTO = SRTT + 4 * RTTVAR$
- Minimum RTO bound = 200ms
  - Actual RTO Timer =  $\max(200ms, RTO)$

# The Incast Workload

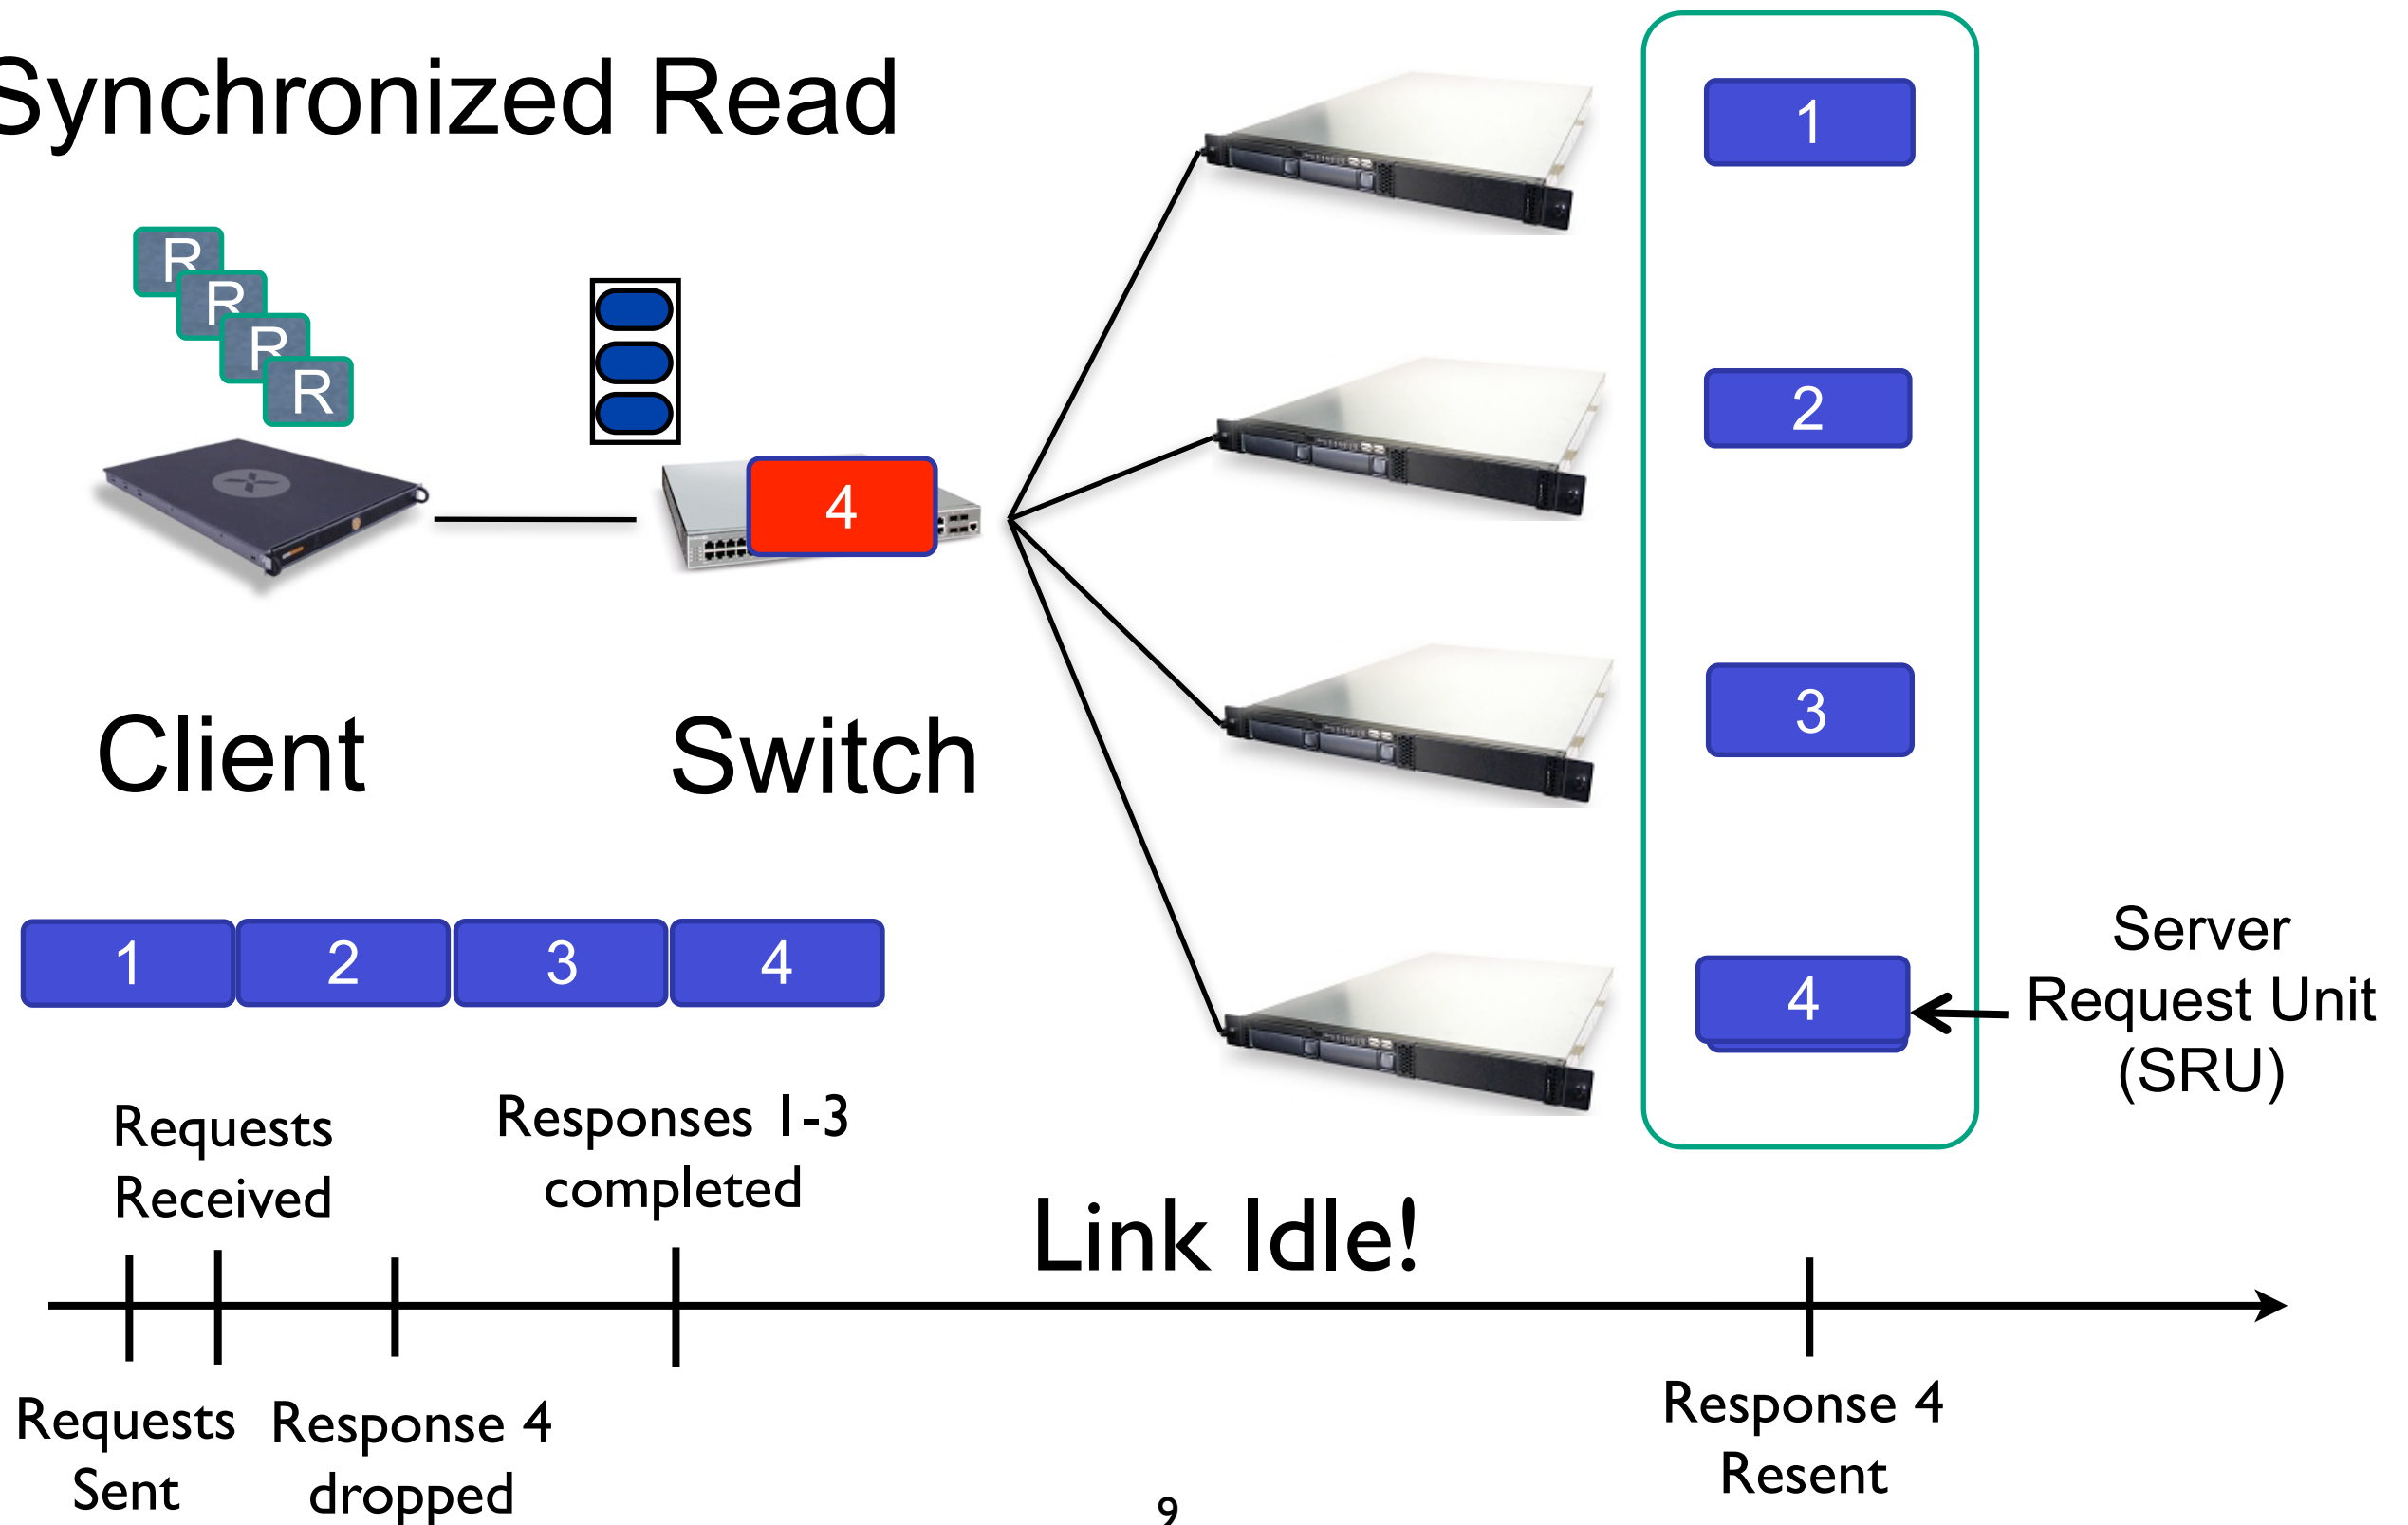
Synchronized Read






# Incast Workload Overfills Buffers


## Synchronized Read



# Client Link Utilization



# 200ms Timeouts Cause Throughput Collapse



## Cluster Environment

1 Gbps Ethernet

100us Delay

200ms RTO

S50 Switch

1 MB Block Size

- [Nagle04] called this Incast; provided app-level workaround
- Cause of throughput collapse: 200ms TCP Timeouts
- Prior work: Other TCP variants did not prevent TCP timeouts. [Phanishayee:FAST2008]

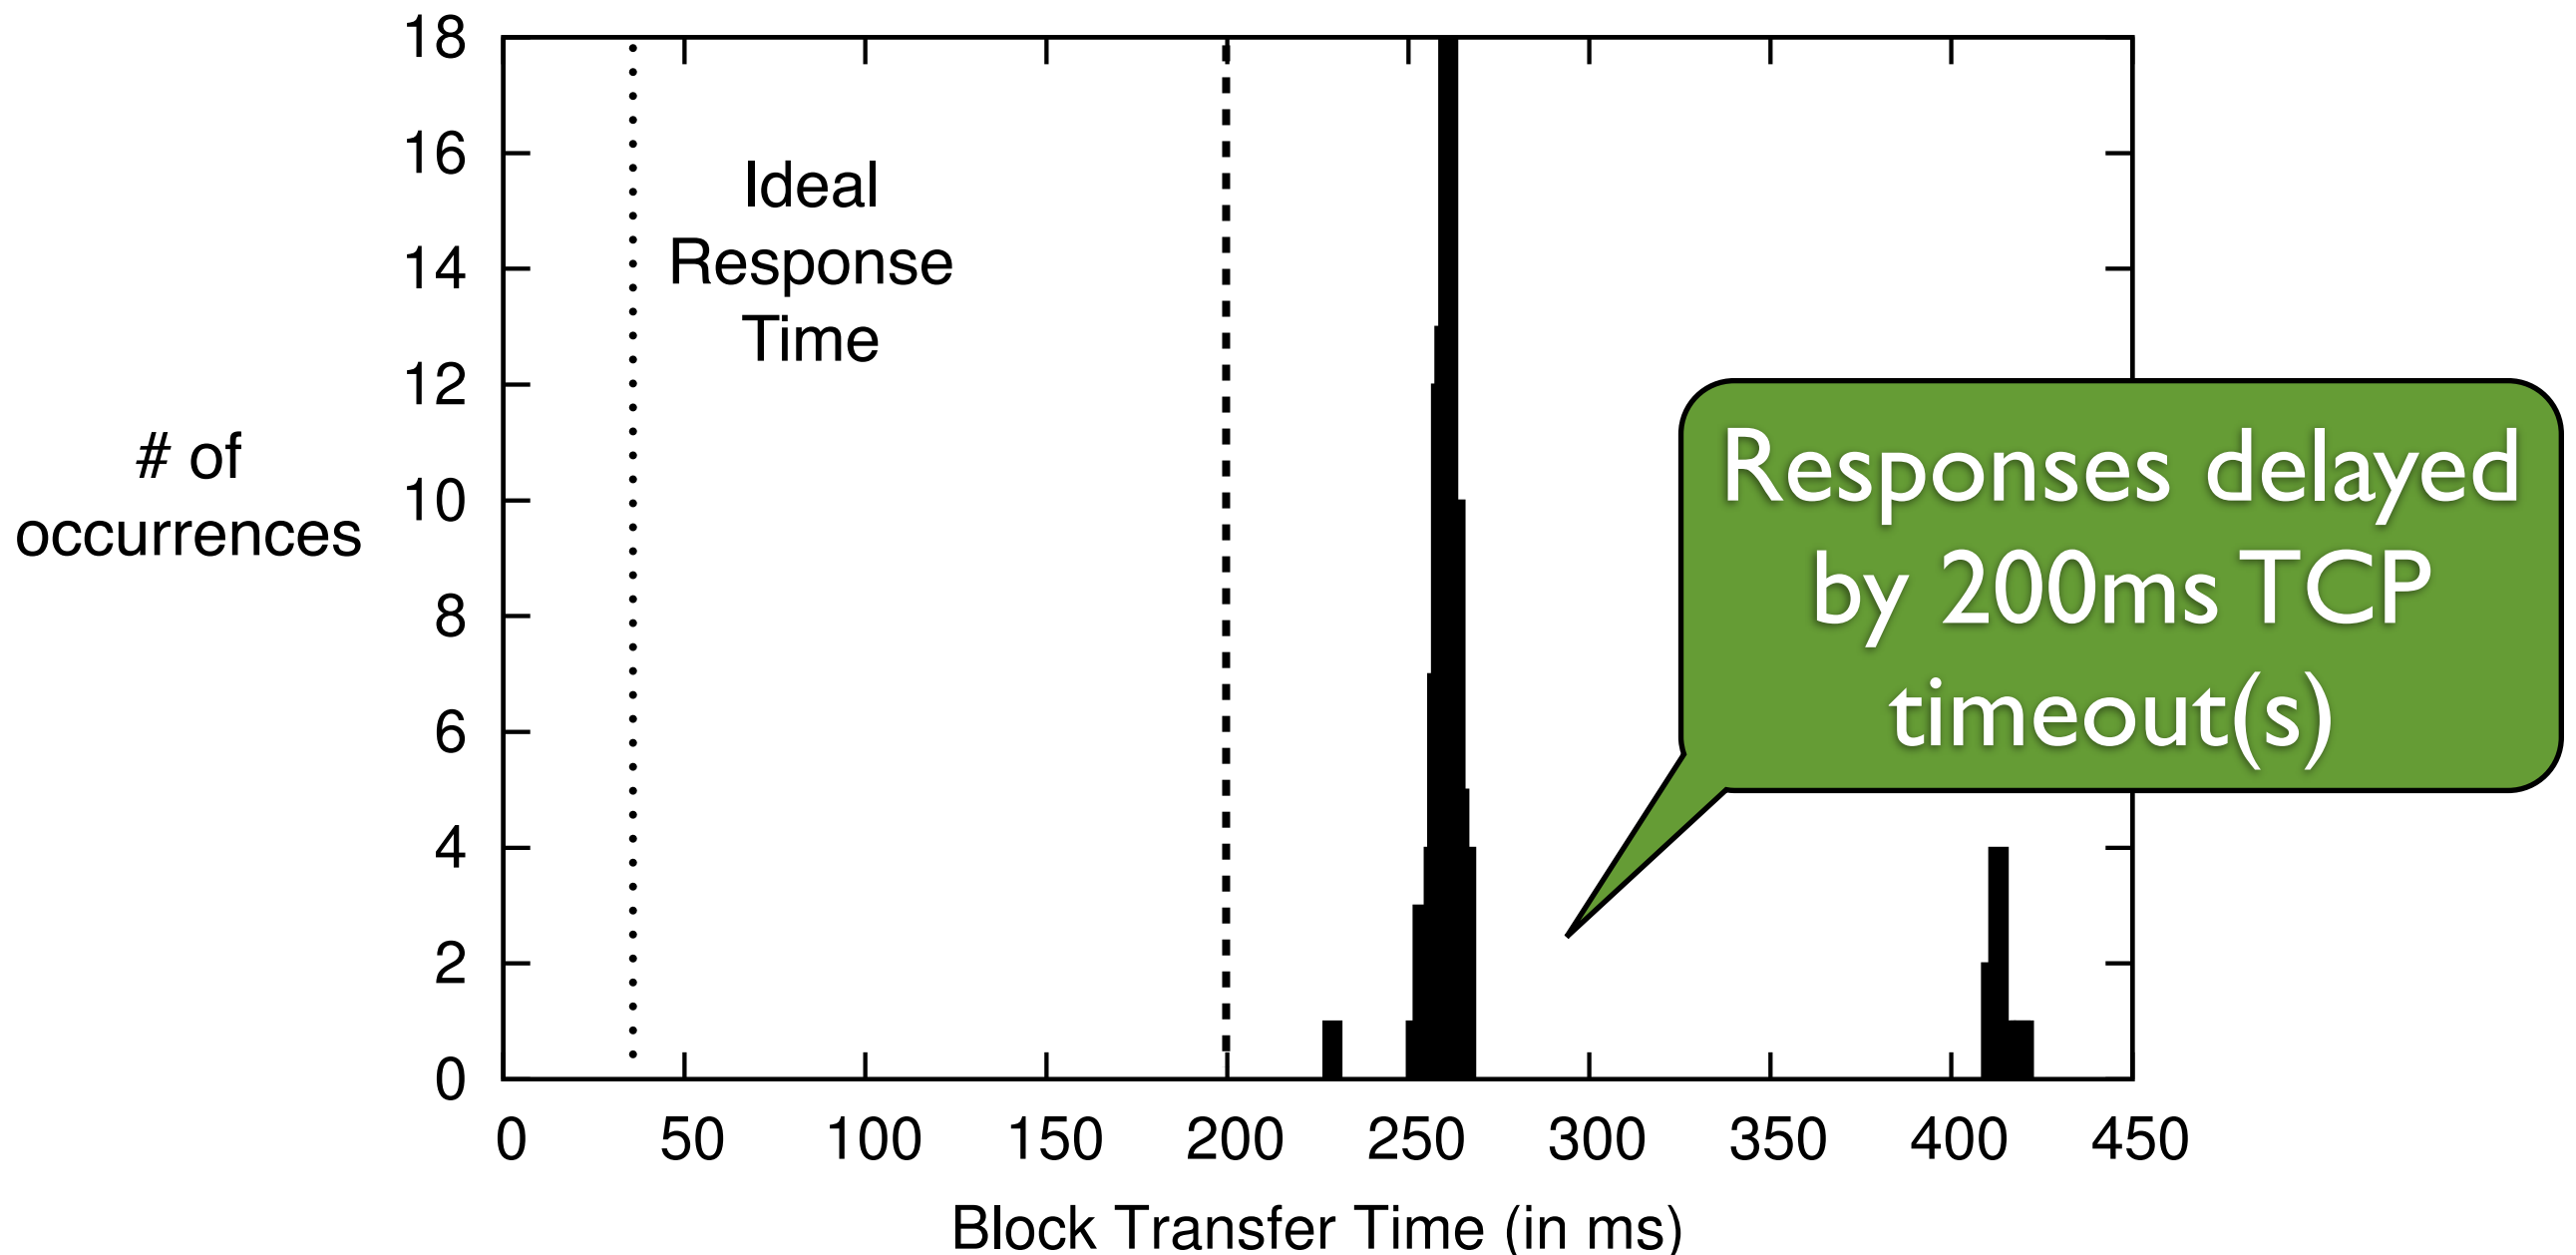
# Latency-sensitive Apps

- Request for 4MB of data sharded across 16 servers (256KB each)
- How long does it take for all of the 4MB of data to return?

# Timeouts Increase Latency

(256KB from 16 servers)

4MB Block Transfer Time Distribution with No RTO bound



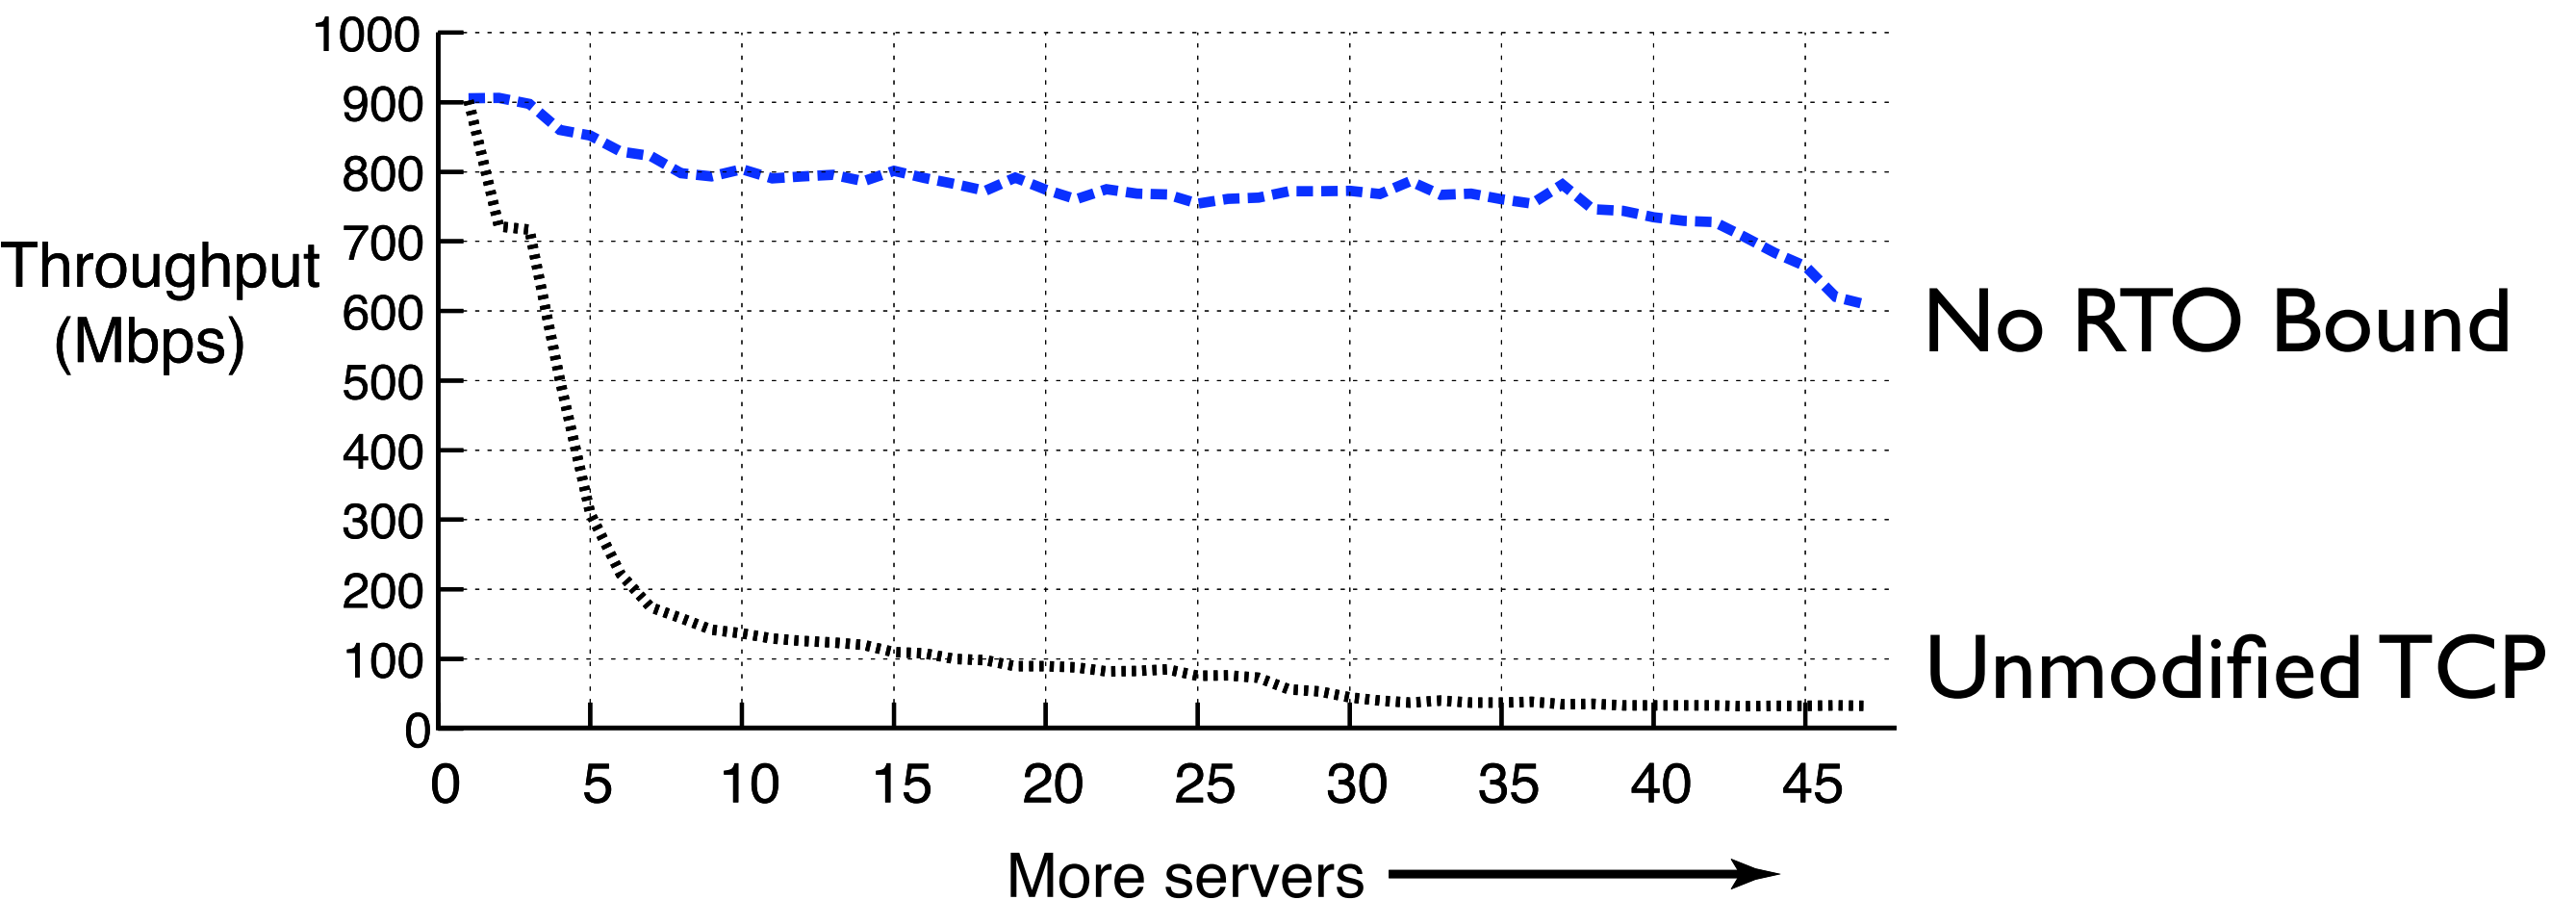
# Outline

- Problem Description, Examples
- **Solution: Microsecond TCP Retransmissions**
- Is it safe?

# Eliminate minRTO

- ✓ Simple one-line change in Linux
  - Does not change RTT measurement granularity
  - Still uses low-resolution, 1ms kernel timers

# Eliminating the RTO bound helps




Millisecond retransmissions not enough



# Requirements for Microsecond RTO

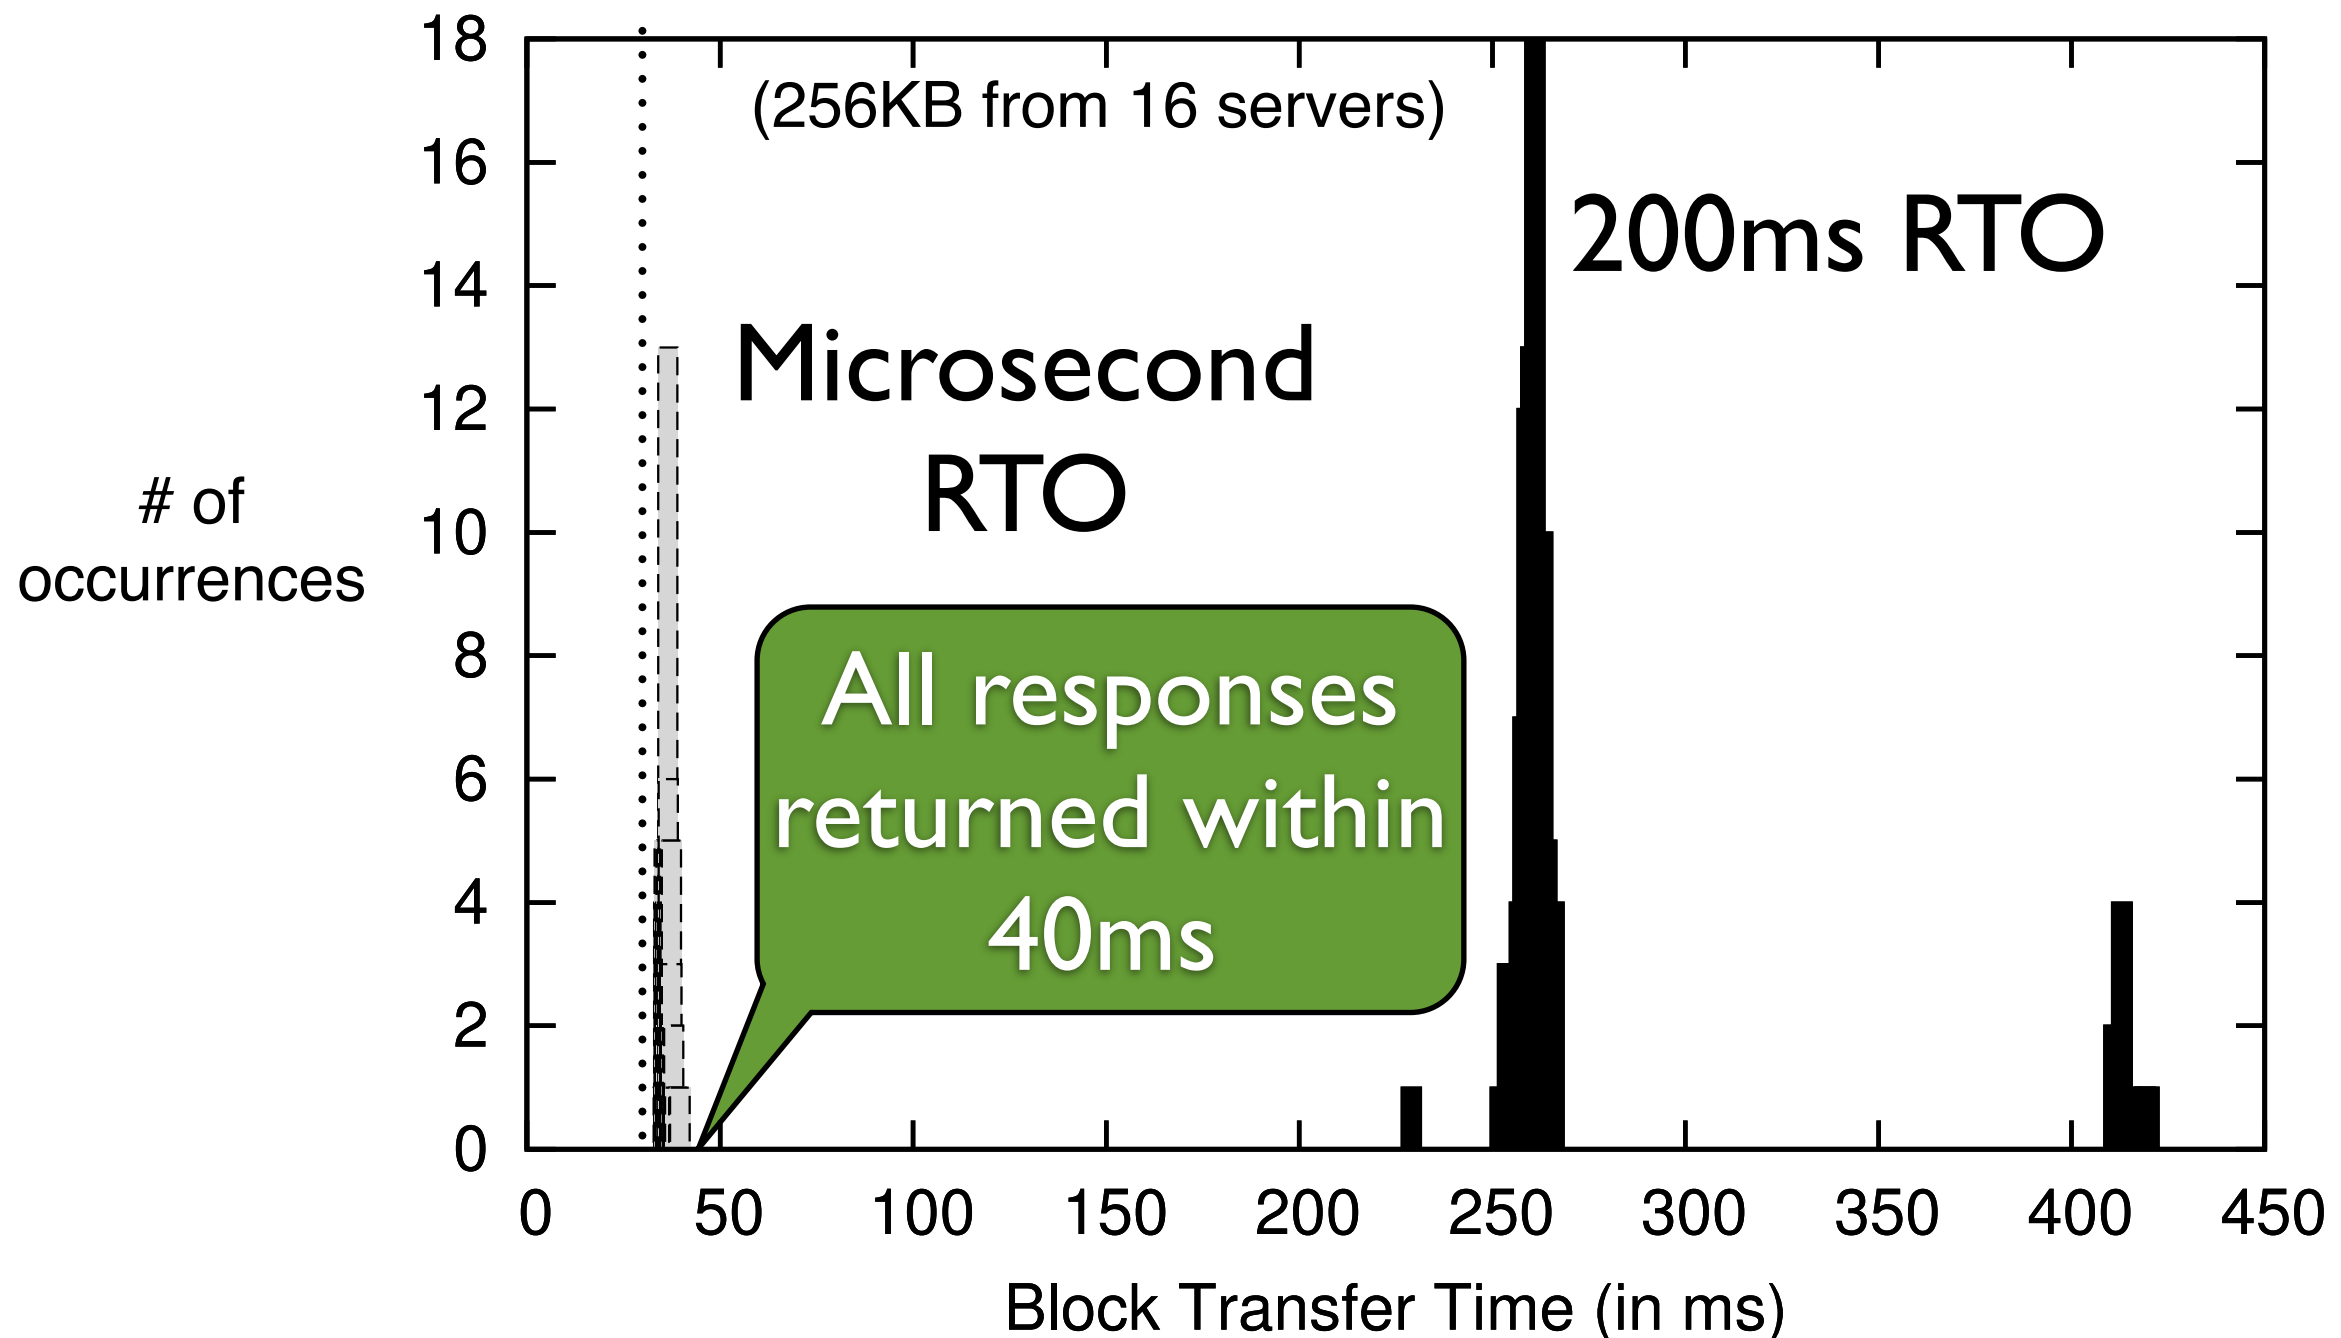
- TCP must track RTT in microseconds
  - Modify internal data-structures
  - Timestamp option (backwards compatible)
- Efficient high-resolution kernel timers
  - Use HPET for efficient interrupt signaling

# Microsecond timeouts are necessary



# Improvement to Latency

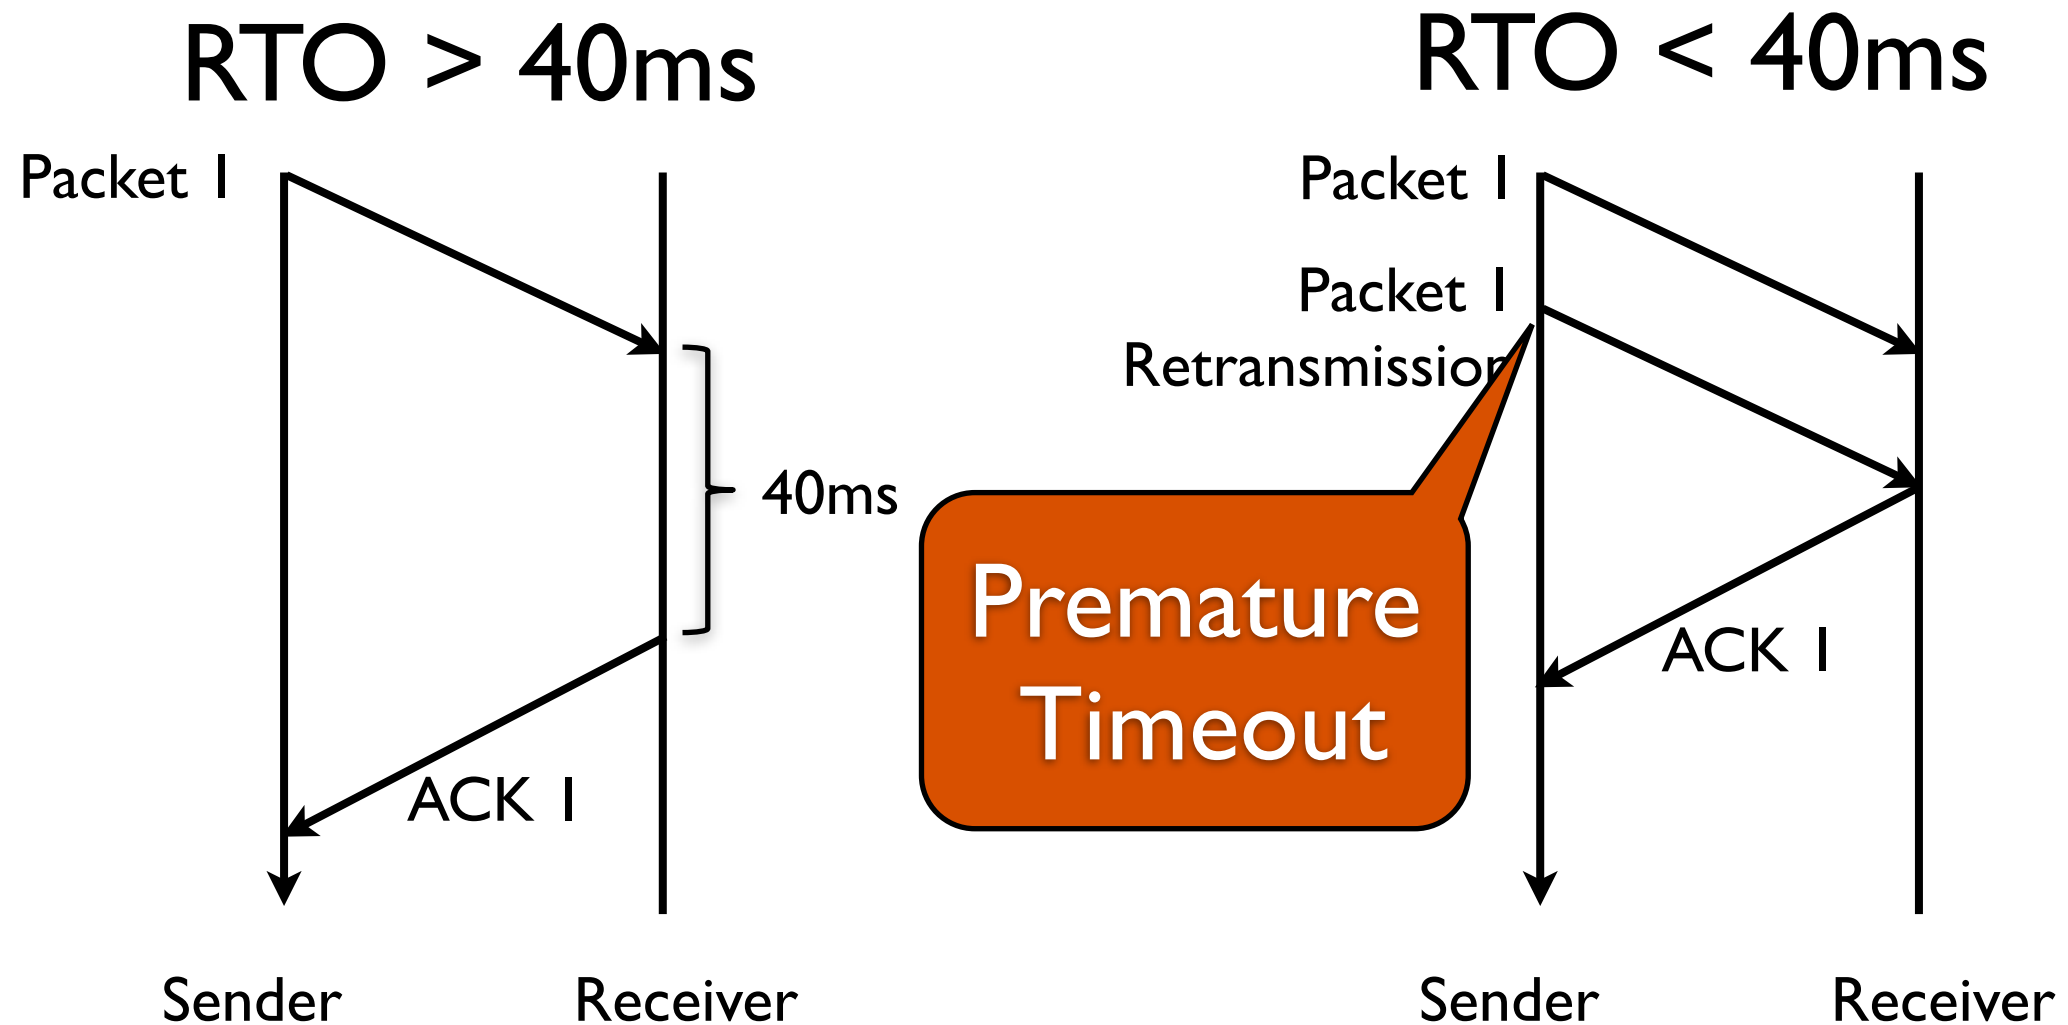
4MB Block Transfer Time Distribution with No RTO bound



# Outline

- Problem Description, Examples
- Solution: Microsecond TCP Retransmissions
- **Is it safe?**
  - Interaction with Delayed ACK
  - Performance in the wide-area


# $\mu$ s RTO and Delayed ACK



RTO on sender triggers  
before delayed ACK on receiver

# Impact of Delayed ACKs

(Fixed Block = 1MB, buffer = 32KB (est.), Switch = Procurve)



10-15% loss in throughput  
with delayed ACK clients

# Is it safe for the wide area?

- Potential Concerns:
  - **Stability**: Could we cause congestion collapse?
  - **Performance**: Do we often timeout unnecessarily?
- Stability preserved
  - Timeouts retain exponential backoff
  - Spurious timeouts *slow* rate of transfer
- Performance: spurious timeouts vs. timely response
  - No optimal RTO estimator [Allman99]

# Spurious timeouts less harmful

Today's TCP has mechanisms to:

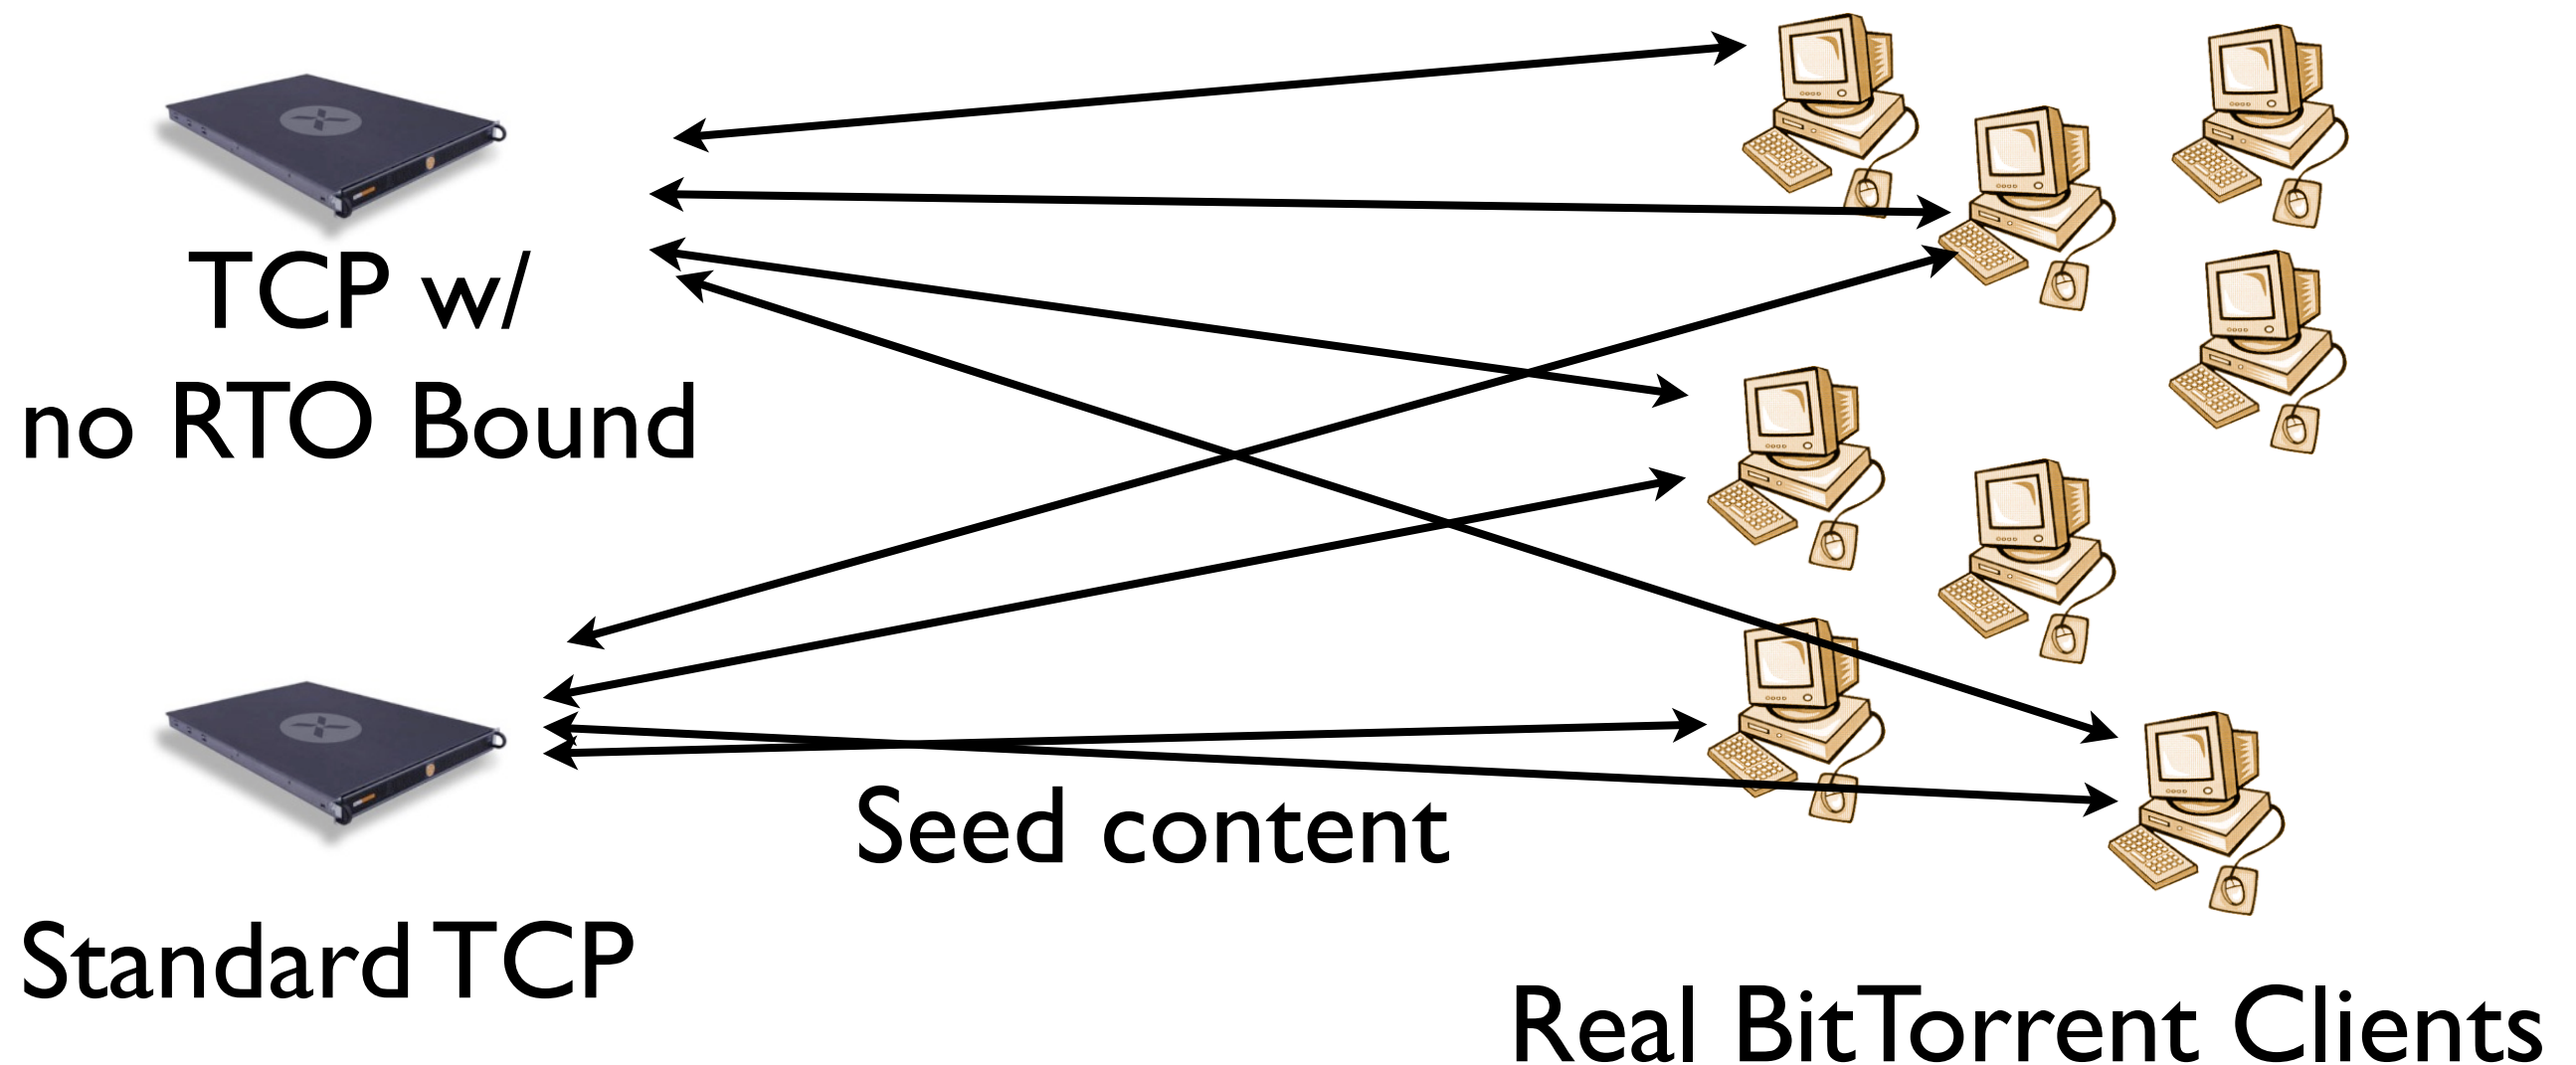
1. Detect spurious timeouts
  - Using TCP timestamp option
2. *Recover* from spurious timeouts
  - Forward RTO (F-RTO)

Both implemented widely!

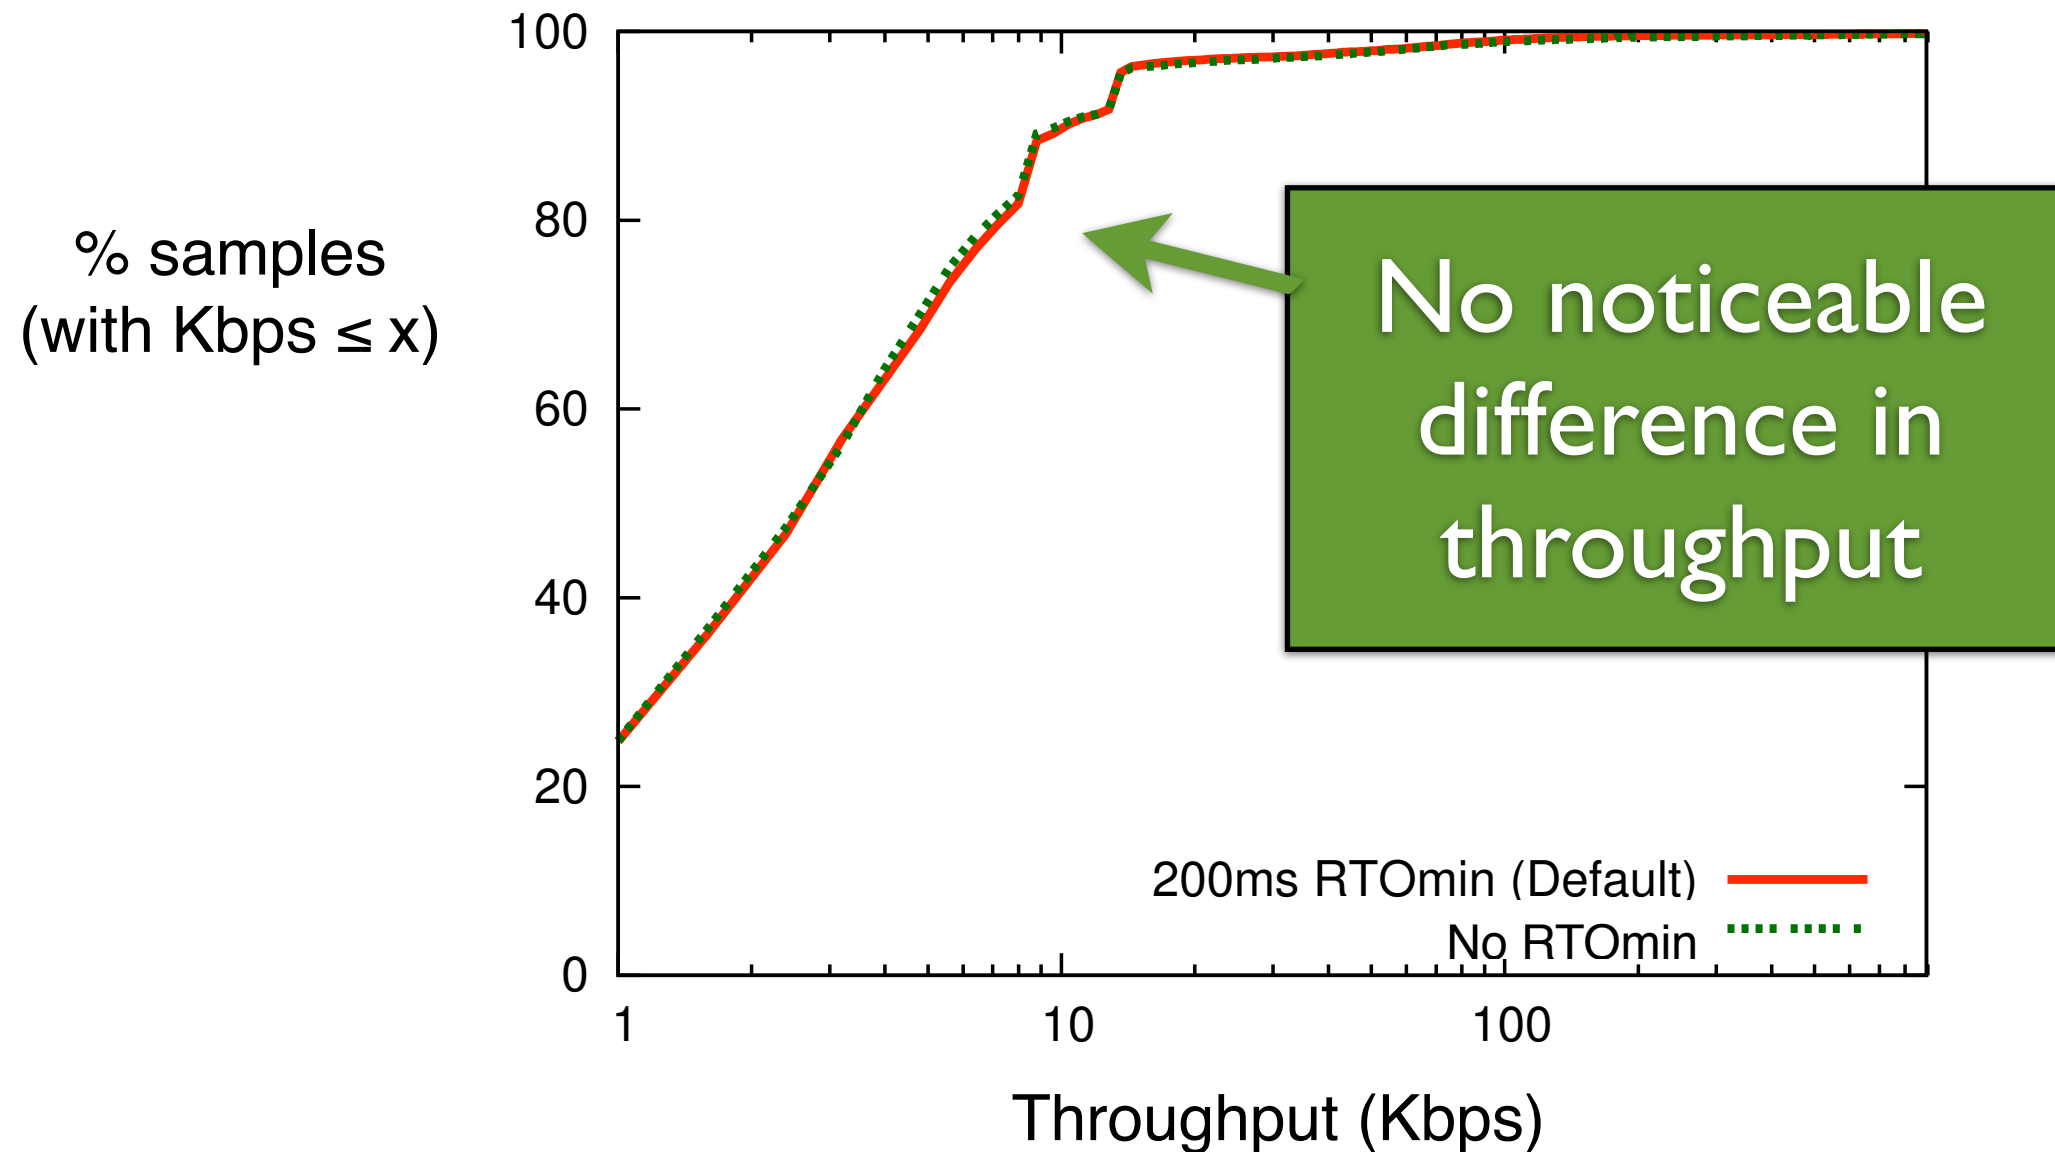


# Wide-area Performance Without minRTO

- Do microsecond timeouts harm wide-area throughput?



# Wide-area Results



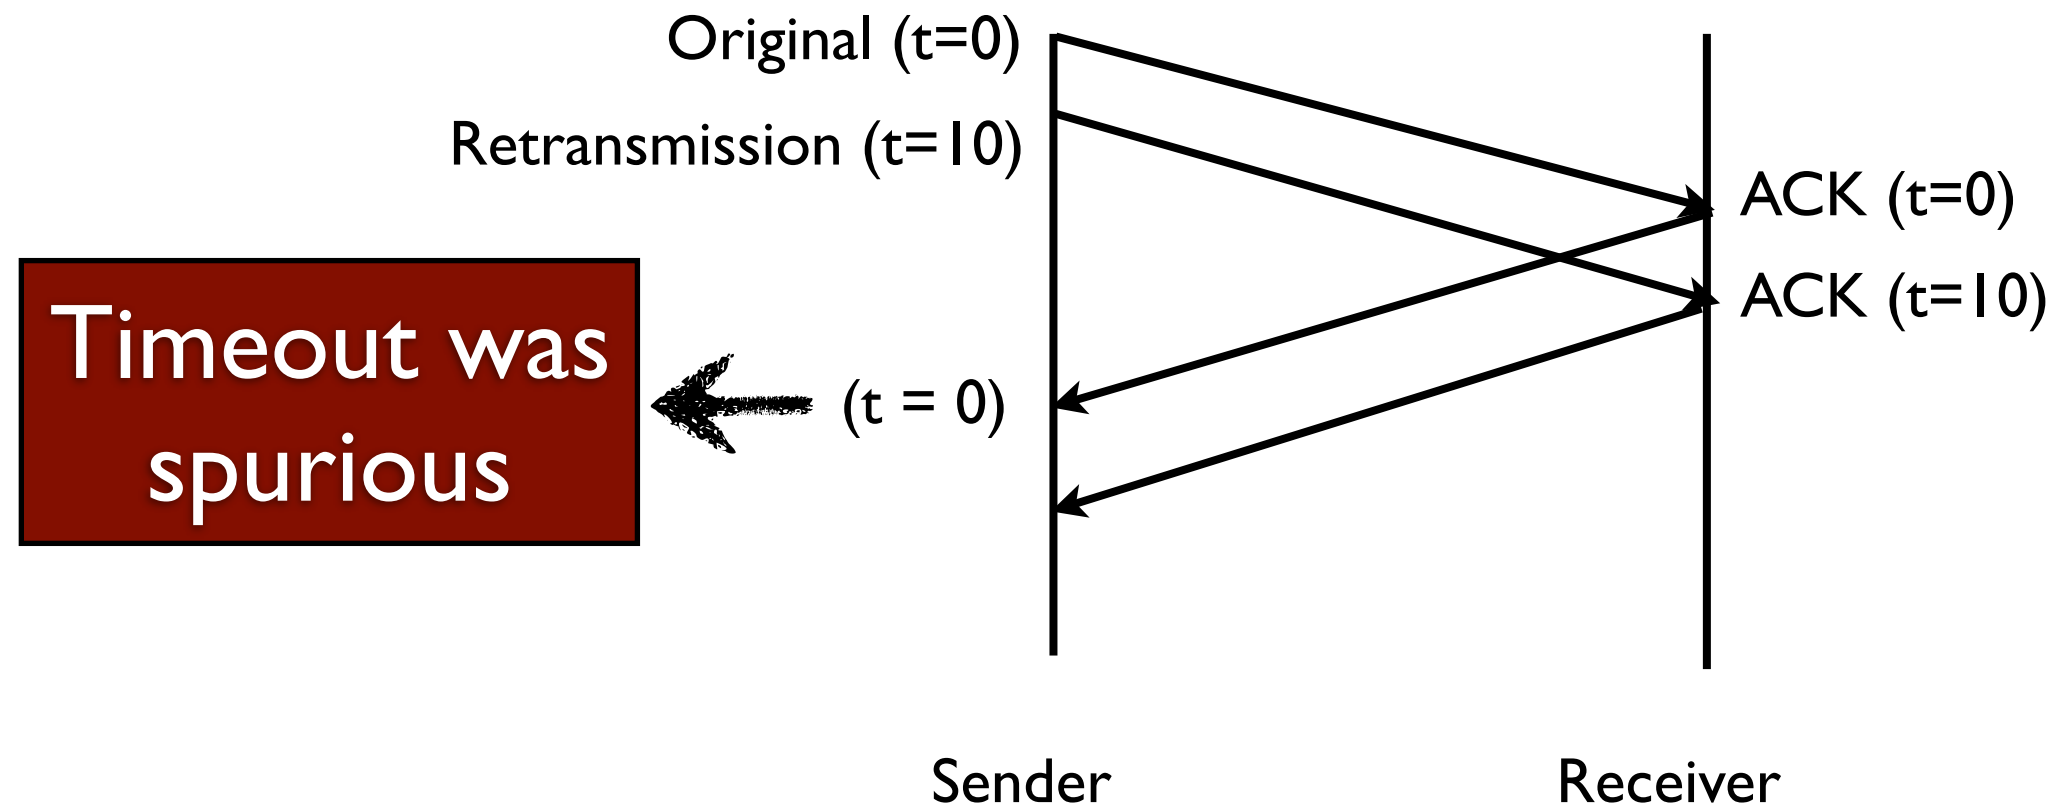
- Few total timeouts (spurious or legitimate)

# Conclusion

- Microsecond RTOs can help datacenter application response time and throughput
- Safe for wide-area communication as well
- Linux patch available:
  - <http://www.cs.cmu.edu/~vrv/incast/>


# Reasons for Relief in 2009

- No ACK ambiguity with TCP timestamp option



- Forward RTO Recovery (RFC4138) more widely deployed (standard in Linux)

# The Need for Microsecond Timeouts



**Simulation Environment**

10Gbps Ethernet

20us Delay

40MB Block Size

- Future datacenters: More bandwidth, less delay, more servers
- Retransmissions should not be bounded

# Control Parameter

